

Camilo Guajardo Inzunza

San Francisco, CA · john.doe@email.com · rendercv.com · rendercv · rendercv

Welcome to RenderCV

RenderCV reads a CV written in a YAML file, and generates a PDF with professional typography.

Each section title is arbitrary.

You can choose any of the 9 entry types for each section.

Markdown syntax is supported everywhere. This is **bold**, *italic*, and [link](#).

Education

Princeton University – Princeton, NJ

Sept 2018 – May 2023

PhD in Computer Science

- ◆ Thesis: Efficient Neural Architecture Search for Resource-Constrained Deployment
- ◆ Advisor: Prof. Sanjeev Arora
- ◆ NSF Graduate Research Fellowship, Siebel Scholar (Class of 2022)

Boğaziçi University – İstanbul, Türkiye

Sept 2014 – June 2018

BS in Computer Engineering

- ◆ GPA: 3.97/4.00, Valedictorian
- ◆ Fulbright Scholarship recipient for Graduate Studies

Experience

Nexus AI – San Francisco, CA

June 2023 – present

Co-Founder & CTO

- ◆ Built foundation model infrastructure serving 2M+ monthly API requests with 99.97% uptime
- ◆ Raised \$18M Series A led by Sequoia Capital, with participation from a16z and Founders Fund
- ◆ Scaled engineering team from 3 to 28 across ML research, platform, and applied AI divisions
- ◆ Developed proprietary inference optimization reducing latency by 73% compared to baseline

NVIDIA Research – Santa Clara, CA

May 2022 – Aug 2022

Research Intern

- ◆ Designed sparse attention mechanism reducing transformer memory footprint by 4.2x
- ◆ Co-authored paper accepted at NeurIPS 2022 (spotlight presentation, top 5% of submissions)

Google DeepMind – London, UK

May 2021 – Aug 2021

Research Intern

- ◆ Developed reinforcement learning algorithms for multi-agent coordination
- ◆ Published research at top-tier venues with significant academic impact
- ◆ ICML 2022 main conference paper, cited 340+ times within two years
- ◆ NeurIPS 2022 workshop paper on emergent communication protocols
- ◆ Invited journal extension in JMLR (2023)

Apple ML Research – Cupertino, CA

May 2020 – Aug 2020

Research Intern

- ◆ Created on-device neural network compression pipeline deployed across 50M+ devices
- ◆ Filed 2 patents on efficient model quantization techniques for edge inference

Microsoft Research – Redmond, WA

May 2019 – Aug 2019

Research Intern

- ◆ Implemented novel self-supervised learning framework for low-resource language modeling
- ◆ Research integrated into Azure Cognitive Services, reducing training data requirements by 60%

Projects

FlashInfer

Jan 2023 – present

Open-source library for high-performance LLM inference kernels

- ◆ Achieved 2.8x speedup over baseline attention implementations on A100 GPUs
- ◆ Adopted by 3 major AI labs, 8,500+ GitHub stars, 200+ contributors

NeuralPrune

Jan 2021

Automated neural network pruning toolkit with differentiable masks

- ◆ Reduced model size by 90% with less than 1% accuracy degradation on ImageNet
- ◆ Featured in PyTorch ecosystem tools, 4,200+ GitHub stars

Publications

Sparse Mixture-of-Experts at Scale: Efficient Routing for Trillion-Parameter Models

July 2023

John Doe, Sarah Williams, David Park

[10.1234/neurips.2023.1234](#) (NeurIPS 2023)

Neural Architecture Search via Differentiable Pruning

Dec 2022

James Liu, John Doe

[10.1234/neurips.2022.5678](#) (NeurIPS 2022, Spotlight)

Multi-Agent Reinforcement Learning with Emergent Communication

July 2022

Maria Garcia, John Doe, Tom Anderson

[10.1234/icml.2022.9012](#) (ICML 2022)

On-Device Model Compression via Learned Quantization

May 2021

John Doe, Kevin Wu

[10.1234/iclr.2021.3456](#) (ICLR 2021, Best Paper Award)

Selected Honors

- ◆ MIT Technology Review 35 Under 35 Innovators (2024)
- ◆ Forbes 30 Under 30 in Enterprise Technology (2024)
- ◆ ACM Doctoral Dissertation Award Honorable Mention (2023)
- ◆ Google PhD Fellowship in Machine Learning (2020 – 2023)
- ◆ Fulbright Scholarship for Graduate Studies (2018)

Skills

Languages: Python, C++, CUDA, Rust, Julia

ML Frameworks: PyTorch, JAX, TensorFlow, Triton, ONNX

Infrastructure: Kubernetes, Ray, distributed training, AWS, GCP

Research Areas: Neural architecture search, model compression, efficient inference, multi-agent RL

Patents

1. Adaptive Quantization for Neural Network Inference on Edge Devices (US Patent 11,234,567)
2. Dynamic Sparsity Patterns for Efficient Transformer Attention (US Patent 11,345,678)
3. Hardware-Aware Neural Architecture Search Method (US Patent 11,456,789)

Invited Talks

4. Scaling Laws for Efficient Inference — Stanford HAI Symposium (2024)
3. Building AI Infrastructure for the Next Decade — TechCrunch Disrupt (2024)
2. From Research to Production: Lessons in ML Systems — NeurIPS Workshop (2023)
1. Efficient Deep Learning: A Practitioner's Perspective — Google Tech Talk (2022)